

Assignment-4 (SMS SPAM Classification)

```
import numpy as np import pandas as pd import seaborn as sns
import matplotlib.pyplot as plt

data =
pd.read_csv('/content/sample_data/spam.csv',delimiter=',',encoding='latin-1') data.head()

      v1                                     v2 Unnamed: 2
\
0  ham  Go until jurong point, crazy.. Available only ...      NaN
1  ham                                     Ok lar... Joking wif u oni...      NaN
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...      NaN
3  ham  U dun say so early hor... U c already then say...      NaN
4  ham  Nah I don't think he goes to usf, he lives aro...      NaN

Unnamed: 3 Unnamed: 4  0
NaN      NaN
1      NaN      NaN
2      NaN      NaN
3      NaN      NaN  4      NaN      NaN data.columns

Index(['v1', 'v2', 'Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],
      dtype='object')

#drop the unnamed columns
data=data.drop(columns=["Unnamed: 2","Unnamed: 3","Unnamed: 4"])

#rename the two relevant columns
data=data.rename(
{
    "v1":"Category",
    "v2":"Message"
},axis=1)
data.head()

      Category                                     Message
0      ham  Go until jurong point, crazy.. Available only ...
1      ham                                     Ok lar... Joking wif u oni...
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...
3      ham  U dun say so early hor... U c already then say...
4      ham  Nah I don't think he goes to usf, he lives aro...
```

```

#check for null values
data.isnull().sum()


Category      0
Message       0
dtype: int64
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
0   Category    5572 non-null   object
1   Message     5572 non-null   object
dtypes: object(2) memory usage: 87.2+
KB

data["Message Length"]=data["Message"].apply(len)

fig=plt.figure(figsize=(12,8))
sns.histplot(
x=data["Message Length"],
hue=data["Category"]
)
plt.title("ham & spam messege length comparision")
plt.show()

```



#Display the description of length of ham and spam messages seperately on an individual series.

```
ham_desc=data[data["Category"]=="ham"]["Message Length"].describe()
spam_desc=data[data["Category"]=="spam"]["Message Length"].describe()
print("Ham Messege Length Description:\n",ham_desc)
print("*****") print("Spam Message
Length Description:\n",spam_desc)
```

Ham Messege Length Description:

```
count    4825.000000
mean      71.023627
std       58.016023
min        2.000000
25%       33.000000
50%       52.000000
75%       92.000000
max      910.000000
```

Name: Message Length, dtype: float64

Spam Message Length Description:

```
count    747.000000 mean
138.866131 std      29.183082
min      13.000000 25%
132.500000
50%      149.000000
```

```
75%      157.000000
max      224.000000
Name: Message Length, dtype: float64
```


```
data.describe(include="all")
```

	Category	Message	Message Length
count	5572	5572	5572.000000
unique	2	5169	NaN
top	ham	Sorry, I'll call later	NaN
freq	4825	30	NaN
mean	NaN	NaN	80.118808
std	NaN	NaN	59.690841
min	NaN	NaN	2.000000
25%	NaN	NaN	36.000000
50%	NaN	NaN	61.000000
75%	NaN	NaN	121.000000
max	NaN	NaN	910.000000

```
data["Category"].value_counts()
```

```
ham      4825
spam     747
Name: Category, dtype: int64
```

```
sns.countplot(
    data=data,
    x="Category"
)
plt.title("ham vs spam")
plt.show()
```



```
ham_count=data["Category"].value_counts()[0]
spam_count=data["Category"].value_counts()[1]
total_count=data.shape[0]

print("Ham contains:{:.2f}% of total
data.".format(ham_count/total_count*100))
print("Spam contains:{:.2f}% of total
data.".format(spam_count/total_count*100))

Ham contains:86.59% of total data.
Spam contains:13.41% of total data.

#compute the length of majority & minority class
minority_len=len(data[data["Category"]=="spam"])
majority_len=len(data[data["Category"]=="ham"])

#store the indices of majority and minority class
minority_indices=data[data["Category"]=="spam"].index
majority_indices=data[data["Category"]=="ham"].index

#generate new majority indices from the total majority_indices
#with size equal to minority class length so we obtain equivalent
number of indices length
random_majority_indices=np.random.choice(
majority_indices,      size=minority_len,
      replace=False
```

```

)

#concatenate the two indices to obtain indices of new dataframe
undersampled_indices=np.concatenate([minority_indices,random_majority_
indices])

#create df using new indices
df=data.loc[undersampled_indices]

#shuffle the sample
df=df.sample(frac=1)

#reset the index as its all mixed
df=df.reset_index()

#drop the older index
df=df.drop(
    columns=["index"],
)

df.shape


(1494, 3)

df["Category"].value_counts()

ham      747
spam      747
Name: Category, dtype: int64

sns.countplot(
    data=df,
    x="Category"
)
plt.title("ham vs spam")
plt.show()

```



```
df.head()
```

	Category	Message	Message
Length			
0	ham	Aah! A cuddle would be lush! I'd need lots of ...	
87			
1	ham	I'm in solihull, do you want anything?	
40			
2	spam	Double Mins & 1000 txts on Orange tariffs. Lat...	
151			
3	ham	No we put party 7 days a week and study lightl...	
126			
4	spam	URGENT!! Your 4* Costa Del Sol Holiday or å£50...	

```
161
```

```
#Created new column Label and encode ham as 0 and spam as 1
```

```
df["Label"]=df["Category"].map(
    {
        "ham":0,
        "spam":1
    }
)
```

```
df.head()
```

	Category	Message	Message
Length \			

```

0      ham  Aah! A cuddle would be lush! I'd need lots of ...
      87
1      ham           I'm in solihull, | do you want anything?
40
2      spam  Double Mins & 1000 txts on Orange tariffs. Lat...
151
3      ham  No we put party 7 days a week and study lightl...
126
4      spam  URGENT!! Your 4* Costa Del Sol Holiday or å£50...
      161

```

```

      Label
0      0
1      0
2      1
3      0  4      1

```

```

import re
import nltk
from nltk.corpus import stopwords

from nltk.stem import PorterStemmer

stemmer=PorterStemmer()

nltk.download('stopwords')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip. True

#declare empty list to store tokenized message
corpus=[]

#iterate through the df["Message"]
for message in df["Message"]:

    #replace every special characters, numbers etc.. with whitespace
of message
    #It will help retain only letter/alphabets
    message=re.sub("[^a-zA-Z]", " ",message)

    #convert every letters to its lowercase
    message=message.lower()

    #split the word into individual word list
    message=message.split()

    #perform stemming using PorterStemmer for all non-englishstopwords
    message=[stemmer.stem(words)
            for words in message

```



```

        if words not in set(stopwords.words("english"))
    ]
    #join the word lists with the whitespace
message=" ".join(message)

    #append the message in corpus list
corpus.append(message)

from tensorflow.keras.preprocessing.text import one_hot
vocab_size=10000


oneHot_doc=[one_hot(words,n=vocab_size)
for words in corpus
]

df["Message Length"].describe()

count      1494.000000
mean        105.203481
std         61.166448
min         3.000000
25%         48.000000
50%        118.000000
75%        153.000000
max         790.000000
Name: Message Length, dtype: float64

fig=plt.figure(figsize=(12,8))
sns.kdeplot(x=df["Message
Length"],
hue=df["Category"]
)
plt.title("ham & spam messege length comparision")
plt.show()

```



```
from tensorflow.keras.preprocessing.sequence import pad_sequences  
sentence_len=200 embedded_doc=pad_sequences(oneHot_doc,  
maxlen=sentence_len,padding="pre")  
  
extract_features=pd.DataFrame(  
    data=embedded_doc  
)  
target=df["Label"]  
  
df_final=pd.concat([extract_features,target],axis=1)  
  
df_final.head()
```

	197	198	199	Label
0	6133	8348	4198	0
1	8663	4425	6636	0
2	4162	8536	7201	1
3	8030	8630	2977	0
4	8861	2485	6055	1

[5 rows x 201 columns]

```
X=df_final.drop("Label",axis=1)
```

```
y=df_final["Label"]
```

```
from sklearn.model_selection import train_test_split
```

```
X_trainval,X_test,y_trainval,y_test=train_test_split(
X,      y,      random_state=42,      test_size=0.15 )
```

```
X_train,X_val,y_train,y_val=train_test_split(
X_trainval,      y_trainval,
random_state=42,      test_size=0.15
)
```

```
from tensorflow.keras.layers import LSTM from
tensorflow.keras.layers import Dense from
tensorflow.keras.layers import Embedding from
tensorflow.keras.models import Sequential
model=Sequential()
```

```
feature_num=100 model.add(
Embedding(
input_dim=vocab_size,
output_dim=feature_num,
input_length=sentence_len
)
) model.add(
LSTM(
units=128
)
)
```

```
model.add(
```

```
Dense(
```

```
units=1,
```

```
      activation="sigmoid"
```

```
    )
)
```

```

from tensorflow.keras.optimizers import Adam
model.compile(optimizer=Adam(
    learning_rate=0.001
),
    loss="binary_crossentropy",
    metrics=["accuracy"]
)

model.fit(
    X_train, y_train,
    validation_data=(
        X_val, y_val
    ), epochs=10 )

Epoch 1/10
34/34 [=====] - 8s 33ms/step - loss: 0.5258 -
accuracy: 0.7653 - val_loss: 0.3215 - val_accuracy: 0.8691
Epoch 2/10
34/34 [=====] - 1s 16ms/step - loss: 0.1718 -
accuracy: 0.9453 - val_loss: 0.1003 - val_accuracy: 0.9738
Epoch 3/10
34/34 [=====] - 1s 16ms/step - loss: 0.0533 -
accuracy: 0.9842 - val_loss: 0.0764 - val_accuracy: 0.9791
Epoch 4/10
34/34 [=====] - 1s 15ms/step - loss: 0.0254 -
accuracy: 0.9926 - val_loss: 0.0716 - val_accuracy: 0.9843
Epoch 5/10
34/34 [=====] - 1s 16ms/step - loss: 0.0184 -
accuracy: 0.9954 - val_loss: 0.0728 - val_accuracy: 0.9843
Epoch 6/10
34/34 [=====] - 1s 16ms/step - loss: 0.0134 -
accuracy: 0.9963 - val_loss: 0.0852 - val_accuracy: 0.9843
Epoch 7/10
34/34 [=====] - 1s 16ms/step - loss: 0.0150 -
accuracy: 0.9954 - val_loss: 0.0744 - val_accuracy: 0.9791
Epoch 8/10
34/34 [=====] - 1s 16ms/step - loss: 0.0112 -
accuracy: 0.9972 - val_loss: 0.0657 - val_accuracy: 0.9843
Epoch 9/10
34/34 [=====] - 1s 16ms/step - loss: 0.0062
accuracy: 0.9981 - val_loss: 0.0732 - val_accuracy: 0.9843
Epoch 10/10
34/34 [=====] - 1s 16ms/step - loss: 0.0050
accuracy: 0.9991 - val_loss: 0.0843 - val_accuracy: 0.9843

<keras.callbacks.History at 0x7fa3263a7850>

y_pred=model.predict(X_test)
y_pred=(y_pred>0.5)

```

8/8 [=====] - 0s 8ms/step from

```
sklearn.metrics import accuracy_score, confusion_matrix
```

```
score=accuracy_score(y_test,y_pred)
```

```
print("Test Score:{:.2f}%".format(score*100))
```

Test Score:96.00%

```
cm=confusion_matrix(y_test,y_pred)
```

```
fig=plt.figure(figsize=(12,8))
```

```
sns.heatmap(      cm,
```

```
annot=True,
```


```
)
```

```
plt.title("Confusion Matrix")
```

```
cm
```

```
array([[100,    2],
```

```
[  7, 116]])
```



#The function take model and message as parameter

```
def classify_message(model,message):
```

#We will treat message as a paragraphs containing multiple sentences(lines)

```

    #we will extract individual lines          for
sentences in message:
sentences=nltk.sent_tokenize(message)

    #Iterate over individual sentences
for sentence in sentences:          #replace
    all special characters
words=re.sub("[^a-zA-Z]", " ",sentence)

    #perform word tokenization of all non-english-stopwords
if words not in set(stopwords.words('english')):
word=nltk.word_tokenize(words)
    word=" ".join(word)

    #perform one_hot on tokenized word
oneHot=[one_hot(word,n=vocab_size)]

    #create an embedded documnet using pad_sequences
    #this can be fed to our model
text=pad_sequences(oneHot,maxlen=sentence_len,padding="pre")
#predict the text using model    predict=model.predict(text)

    #if predict value is greater than 0.5 its a spam
if predict>0.5:    print("It is a spam")
    #else the message is not a spam
else:    print("It is not a
spam")

message1="I am having a bad day and I would like to have a break
today"
message2="This is to inform you had won a lottery and the subscription
will end in a week so call us." nltk.download('punkt')

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip. True

classify_message(model,message1)

1/1 [=====] - 0s 21ms/step
It is not a spam

classify_message(model,message2)

1/1 [=====] - 0s 22ms/step
It is a spam

```