

```


import cv2
import numpy as np
from keras.datasets import mnist
from keras.layers import Dense, Flatten, MaxPooling2D, Dropout
from keras.layers.convolutional import Conv2D
from keras.models import Sequential
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

(X_train, y_train), (X_test, y_test) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step

plt.imshow(X_train[0], cmap="gray")
plt.show()
print (y_train[0])

```



5

```

print ("Shape of X_train: {}".format(X_train.shape))
print ("Shape of y_train: {}".format(y_train.shape))
print ("Shape of X_test: {}".format(X_test.shape))
print ("Shape of y_test: {}".format(y_test.shape))

Shape of X_train: (60000, 28, 28)
Shape of y_train: (60000,)
Shape of X_test: (10000, 28, 28)
Shape of y_test: (10000,)

```

```

# Reshaping so as to convert images for our model
X_train = X_train.reshape(60000, 28, 28, 1)
X_test = X_test.reshape(10000, 28, 28, 1)

print ("Shape of X_train: {}".format(X_train.shape))
print ("Shape of y_train: {}".format(y_train.shape))
print ("Shape of X_test: {}".format(X_test.shape))
print ("Shape of y_test: {}".format(y_test.shape))

Shape of X_train: (60000, 28, 28, 1)
Shape of y_train: (60000,)
Shape of X_test: (10000, 28, 28, 1)
Shape of y_test: (10000,)

#one hot encoding
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

model = Sequential()

## Declare the layers
layer_1 = Conv2D(64, kernel_size=3, activation='relu',
input_shape=(28, 28, 1))
layer_2 = MaxPooling2D(pool_size=2)
layer_3 = Conv2D(32, kernel_size=3, activation='relu')
layer_4 = MaxPooling2D(pool_size=2)
layer_5 = Dropout(0.5)
layer_6 = Flatten()
layer_7 = Dense(128, activation="relu")
layer_8 = Dropout(0.5)
layer_9 = Dense(10, activation='softmax')

## Add the layers to the model
model.add(layer_1)
model.add(layer_2)
model.add(layer_3)
model.add(layer_4)
model.add(layer_5)
model.add(layer_6)
model.add(layer_7)
model.add(layer_8)
model.add(layer_9)

model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

model.fit(X_train, y_train, validation_data=(X_test, y_test),
epochs=3)

Epoch 1/3
1875/1875 [=====] - 58s 31ms/step - loss:

```

0.8654 - accuracy: 0.7801 - val_loss: 0.1307 - val_accuracy: 0.9630
Epoch 2/3

1875/1875 [=====] - 58s 31ms/step - loss:
0.2703 - accuracy: 0.9201 - val_loss: 0.0750 - val_accuracy: 0.9757

Epoch 3/3

1875/1875 [=====] - 56s 30ms/step - loss:
0.2055 - accuracy: 0.9385 - val_loss: 0.0746 - val_accuracy: 0.9772

<keras.callbacks.History at 0x7fc1e21cc510>

```
example = X_train[1]
prediction = model.predict(example.reshape(1, 28, 28, 1))
print ("Prediction (Softmax) from the neural network:\n\n
{}".format(prediction))
hard_maxed_prediction = np.zeros(prediction.shape)
hard_maxed_prediction[0][np.argmax(prediction)] = 1
print ("\n\nHard-maxed form of the prediction: \n\n
{}".format(hard_maxed_prediction))
```

```
print ("\n\n----- Prediction ----- \n\n")
plt.imshow(example.reshape(28, 28), cmap="gray")
plt.show()
print("\n\nFinal Output: {}".format(np.argmax(prediction)))
```

1/1 [=====] - 0s 83ms/step


Prediction (Softmax) from the neural network:

```
[[9.99999881e-01 7.21094625e-13 7.90088137e-08 3.49195464e-11
 1.54954244e-11 5.48896974e-13 1.05098525e-08 1.00683108e-10
 7.00186797e-10 1.28125794e-08]]
```

Hard-maxed form of the prediction:

```
[[1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]
```

----- Prediction -----



Final Output: 0

```
metrics=model.evaluate(X_test,y_test,verbose=0)
print("Metrics(test loss and Test Accuracy):")
print(metrics)
```

```
Metrics(test loss and Test Accuracy):
[0.07461030036211014, 0.9771999716758728]
```

```
image = cv2.imread('test_image.jpg')
image = np.full((100,80,3), 12, dtype = np.uint8)
grey = cv2.cvtColor(image.copy(), cv2.COLOR_BGR2GRAY)
ret, thresh = cv2.threshold(grey.copy(), 75, 255,
cv2.THRESH_BINARY_INV)
contours,hierarchy = cv2.findContours(thresh.copy(),
cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
preprocessed_digits = []
```

```
for c in contours:
    x,y,w,h = cv2.boundingRect(c)
```

```
    # Creating a rectangle around the digit in the original image (for
displaying the digits fetched via contours)
    cv2.rectangle(image, (x,y), (x+w, y+h), color=(0, 255, 0),
thickness=2)
```

```
    # Cropping out the digit from the image corresponding to the
current contours in the for loop
```

```

digit = thresh[y:y+h, x:x+w]

# Resizing that digit to (18, 18)
resized_digit = cv2.resize(digit, (18,18))

# Padding the digit with 5 pixels of black color (zeros) in each
side to finally produce the image of (28, 28)
padded_digit = np.pad(resized_digit, ((5,5),(5,5)), "constant",
constant_values=0)

# Adding the preprocessed digit to the list of preprocessed digits
preprocessed_digits.append(padded_digit)

print("\n\n\n-----Contoured Image-----")
import os, types
import pandas as pd

def __iter__(self): return 0

print("\n\n\n-----Contoured Image-----")
plt.imshow(image, cmap="gray")
plt.show()

inp = np.array(preprocessed_digits)

```

-----Contoured Image-----

